

Basics of Parallel Computing
2024S
Assignment 1

May 2, 2024

2 Person Group 13

1: Pia SCHWARZINGER, 12017370

2: Yahya JABARY, 11912007

2 The Tasks

2.2 Compute Speed-up and Parallel Efficiency for 2 Instance Sizes

Tables 1 and 2 present the parallel runtime measurements and the respective speed-up and efficiency values for the cases c_s and c_b .

Table 1: Runtime and speed-up of parallel Julia set generator for c_s case.

size	p	mean runtime (s)	speed-up	par. eff.
155	1	0.276419	1	1.42473
155	2	0.14614	1.89147	1.34742
155	4	0.0850049	3.2518	1.15824
155	8	0.0684728	4.03691	0.71894
155	16	0.0727858	3.7977	0.338169
155	24	0.0845242	3.27029	0.194137
155	32	0.0975262	2.8343	0.126191
1100	1	13.42	1	1.40137
1100	2	6.66189	2.01444	1.41149
1100	4	3.37751	3.97333	1.39203
1100	8	1.72664	7.7723	1.36149
1100	16	0.908652	14.7691	1.29356
1100	24	0.654567	20.5021	1.19713
1100	32	0.564249	23.7838	1.04156

Keep in mind that while the speed-up was calculated using $p = 1$ as the reference point, the parallel efficiency was calculated using an average of the sequential runtime, by running the following commands 3 times on the Hydra-cluster and averaging the results:

```

srun -p q_student -t 1 -N 1 -c 32 python3 julia.py --size 155 --nprocs 1
# 155;20;1;0.39382300106808543
srun -p q_student -t 1 -N 1 -c 32 python3 julia.py --size 1100 --nprocs 1
# 1100;20;1;18.806384983938187

```

Table 2: Runtime and speed-up of parallel Julia set generator for c_b case.

size	p	mean runtime (s)	speed-up	par. eff.
155	1	0.394086	1	0.711286
155	2	0.210897	1.86862	0.66456
155	4	0.121818	3.23504	0.575259
155	8	0.0862224	4.57057	0.406373
155	16	0.0852527	4.62256	0.205497
155	24	0.0963721	4.08921	0.121191
155	32	0.108531	3.6311	0.0807109
1100	1	19.1049	1	0.68691
1100	2	9.67163	1.97536	0.678447
1100	4	4.90018	3.89883	0.669536
1100	8	2.4526	7.78967	0.66885
1100	16	1.28631	14.8525	0.637646
1100	24	0.900398	21.2183	0.607295
1100	32	0.746145	25.6049	0.549633

Same as before, the parallel efficiency was calculated using an average of the sequential runtime, by running the following commands 3 times on the Hydra-cluster and averaging the results:

```

srun -p q_student -t 1 -N 1 -c 32 python3 julia.py --size 155 --nprocs 1 --benchmark
# 155;20;1;0.2803074959665537
srun -p q_student -t 1 -N 1 -c 32 python3 julia.py --size 1100 --nprocs 1 --benchmark
# 1100;20;1;13.123375411145389

```

It's important to mention that we opted for logarithmic scaling on all y-axes in our graphs. This choice was made to enhance the visibility of function shapes/gradients over the marginal differences. However, it's crucial to be mindful that this results in each increment on the y-axis carrying significantly more weight.

Also: Going forward, we won't focus as much on how different seeds (S-Case vs. B-Case) change the Julia set fractal. This is because the load size has a much bigger impact on our results. The results we've seen using different seeds have been very similar overall.

2.2.1 Comparing: Absolute Speed-up vs. Number of Processes

Figure 1 shows the absolute runtime in comparison to number of cores for both the S-case and B-case for two different sizes.

When comparing the impact of process count on runtime across datasets of differing sizes, we observe a few key trends.

With the larger dataset, there's a significantly sharper decline in runtime as we increase the number of processes/processors (these terms can *approximately* be used interchangeably since setting `chunksize=1` in Python's `multiprocessing.Pool.map` means one task is mapped to each processor at a time - unless we have fewer cores than processes). Additionally, the runtime appears to flatten out as we increase the processor count. Conversely, with the smaller dataset, the decline in runtime is less pronounced, and starting from 8 processes, the overhead from fork syscalls actually outweighs the benefits of parallel computation, leading to an increase in runtime.

2.2.2 Comparing: Relative Speed-up vs. Number of Processes

Figure 2 demonstrates the relative speed up in comparison to number of cores for both the S-case and B-case for the same sizes as before.

In our earlier analysis, we delved into the marginal decrease in runtime with the addition of each process. Essentially, this mirrors the concept of speedup, particularly when our baseline is at $p = 1$. These findings reinforce what we've already discussed. However, it's important to note that the speedup gap between the two loads in the S-case is significantly larger. This likely stems from the increased number of iterations and computations needed per pixel to achieve convergence in the visualization.

2.2.3 Comparing: Parallel Efficiency vs. Number of Processes

The final figure 3 reveals a generally similar downward trend for both seed-cases in the significance per process with increasing number of processes (note that this is parallel efficiency, distinct from utilization). Additionally the decrease in parallel efficiency is a lot more pronounced in the smaller workload and comes very close to 0 as we reach $p = 32$ in both seed-cases. A crucial distinction lies in the initial baseline: in the B-case, the starting point for process efficiency is substantially lower (by a factor of 2) than the S-case. This significant difference in initial efficiency is not directly apparent from the graph itself and would require further examination of the data provided in the appendix. This could be another indicator for a higher compute-intensity in the S-case in addition to the gap previously observed in the speed-up.

2.2.4 Discussion

The findings highlight the importance of considering load size and parallel processing efficiency in computational analysis. Variation in seed choice has minimal impact compared to load size on results. Larger datasets show sharper decline in runtime with increasing processes, while smaller ones exhibit diminishing returns and eventual increase in runtime due to overhead. Parallel efficiency decreases with more processes, with the B-case starting at a lower efficiency than the S-case, indicating higher compute-intensity in the latter.

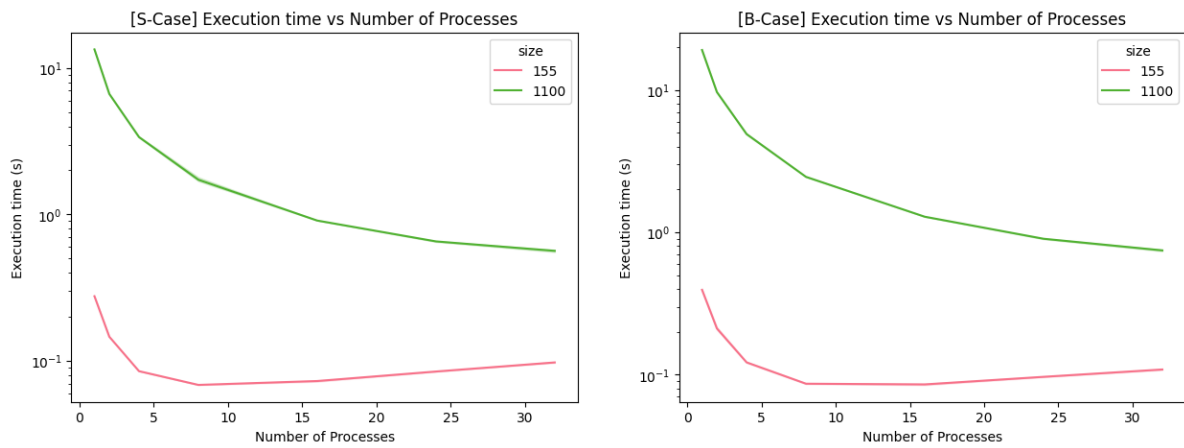


Figure 1: Absolute Runtime vs. Number of Processes

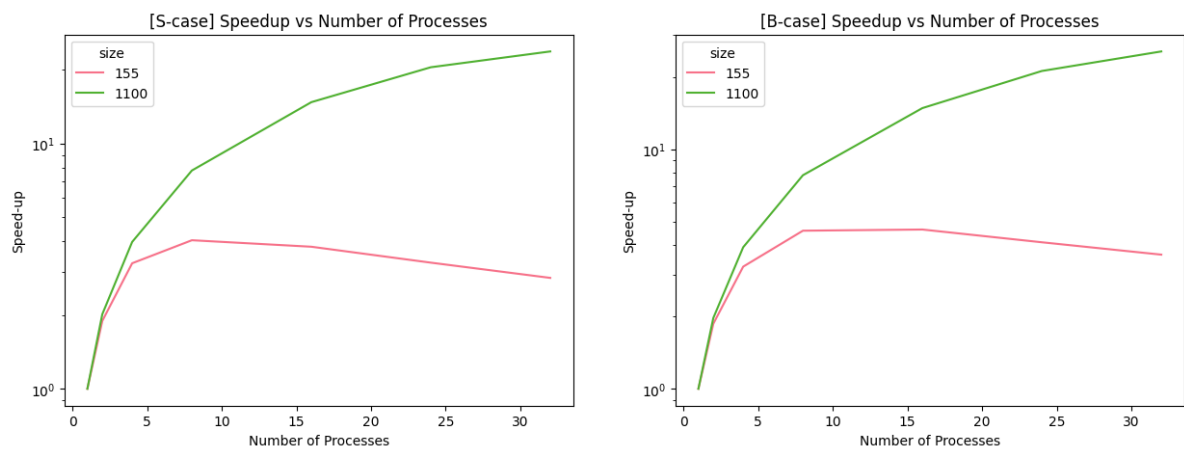


Figure 2: Relative Speed-up vs. Number of Processes

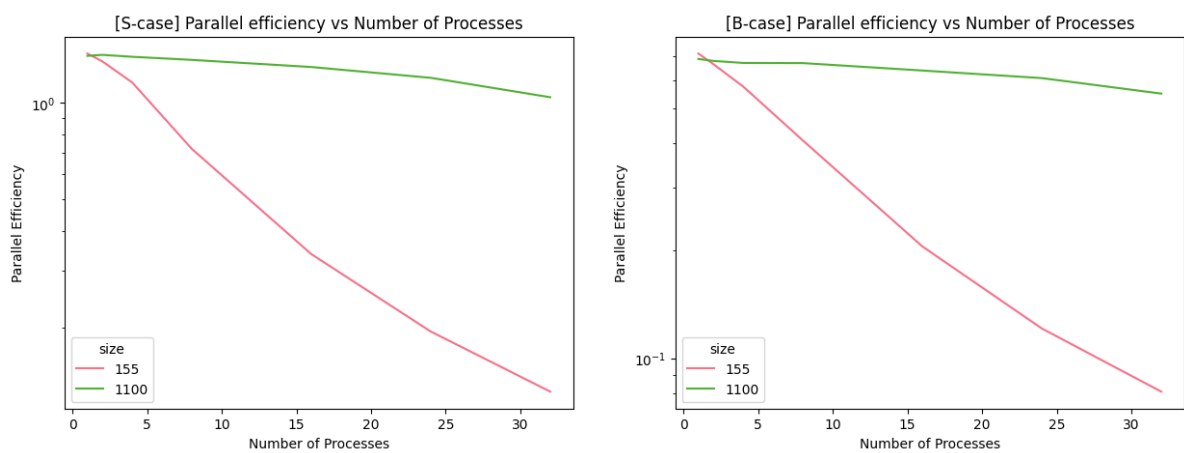


Figure 3: Parallel Efficiency vs. Number of Processes

2.3 Influence of Patch Size

Table 3 presents the average runtimes in seconds of varying patch sizes, while keeping the number of cores and problem size the same. Again, the means are taken from three runs on the cluster.

Table 3: Influence of patch size on mean runtime

size	patch	nprocs	avg_time
950	1	32	59.680444
950	5	32	2.439178
950	10	32	0.827185
950	20	32	0.589660
950	55	32	0.649828
950	150	32	1.791236
950	400	32	5.347212

Figure 4 now visually shows the influence of the patch size on the mean runtime.

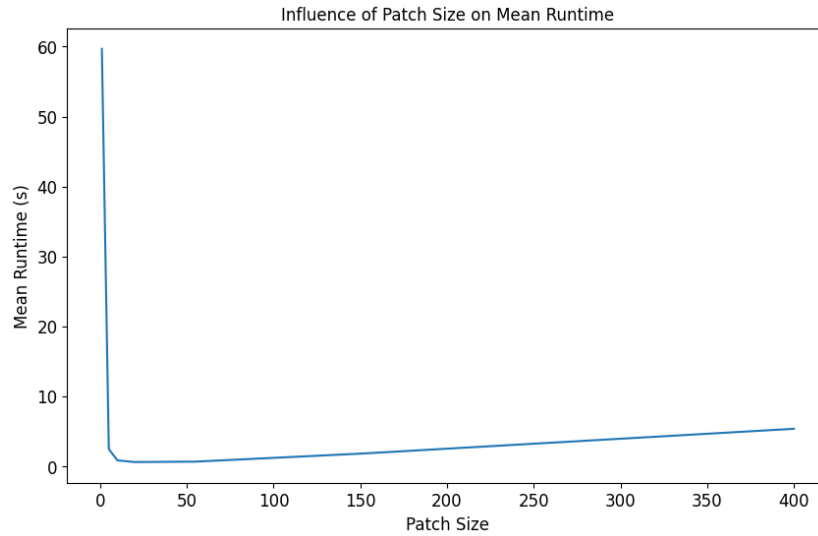


Figure 4: Mean Runtime vs. Patch Size - Influence of patch size

As we can tell having a patch size of 1 leads to an extremely high runtime but drastically decreases once multiple patches are involved. This is because computational resources can be used more effectively and parallel processing can be improved. However, starting from 150 patches upwards, the runtime increases which is probably due to the overhead that comes with managing a large number of parallel processes.

2.4 Finding the Best Patch Size

In order to find the best possible patch size, table 4 and figure 5 present the results of our experiments.

Table 4: Best patch size

size	patch	nprocs	avg_time
800	1	16	50.020686
800	2	16	12.428263
800	3	16	5.235198
800	4	16	2.781848
800	5	16	1.778088
800	6	16	1.262205
800	7	16	1.025805
800	8	16	0.871605
800	9	16	0.841269
800	10	16	0.788733
800	11	16	0.777117
800	12	16	0.730836
800	13	16	0.721792
800	14	16	0.732972
800	15	16	0.712528
800	16	16	0.736369
800	17	16	0.714250
800	18	16	0.709131
800	19	16	0.711833
800	20	16	0.704302
800	21	16	0.712072
800	22	16	0.715015
800	23	16	0.718865
800	24	16	0.710345
800	25	16	0.706855
800	26	16	0.709854
800	27	16	0.715558
800	28	16	0.708274
800	29	16	0.702031
800	30	16	0.705310

Based on figure 5, one can tell that the runtime seems to converge starting from a patch size of 10. The minimum and therefore best patch size according to the numbers provided in table 4, is 29 which leads to an average runtime of 0.702 seconds.

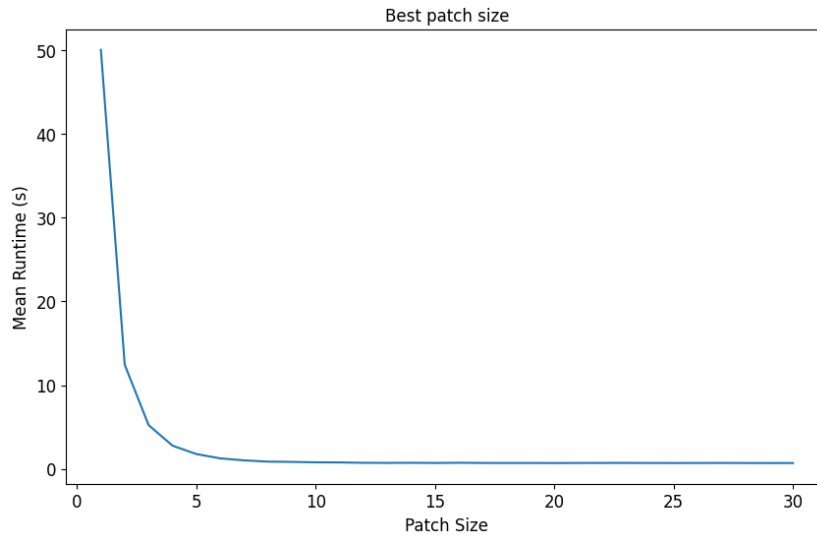


Figure 5: Mean Runtime vs. Patch Size - Finding best patch size

3 Speed-up Analysis

3.1 Absolute speed up

One can calculate the absolute speed up by computing $S_a^{A_i}(n, p) = \frac{T_{seq}^*(n)}{T_{par}(n, p)}$.

Algorithm sequential:

$$T_{seq}^*(n) = O(n \log n)$$

$$T_{seq}^*(1000) = O(1000 \log 1000) \approx 6907.76$$

Algorithm A1:

$$T_{par}^{A_1}(n, p) = O\left(\frac{n \log n}{p} + \log n\right)$$

$$T_{par}^{A_1}(1000, 4) \approx 1733.85$$

$$T_{par}^{A_1}(1000, 16) \approx 438.64$$

$$T_{par}^{A_1}(1000, 64) \approx 114.84$$

$$S_a^{A_1}(1000, 4) = \frac{6907.76}{1733.85} \approx 3.98$$

$$S_a^{A_1}(1000, 16) = \frac{6907.76}{438.64} \approx 15.74$$

$$S_a^{A_1}(1000, 64) = \frac{6907.76}{114.84} \approx 60.15$$

Algorithm A2:

$$T_{\text{par}}^{A_2}(n, p) = O\left(\frac{n \log n}{p} + n\right)$$

$$T_{\text{par}}^{A_2}(1000, 4) \approx 2726.94$$

$$T_{\text{par}}^{A_2}(1000, 16) \approx 1431.73$$

$$T_{\text{par}}^{A_2}(1000, 64) \approx 1107.93$$

$$S_a^{A_2}(1000, 4) = \frac{6907.76}{2726.94} \approx 2.53$$

$$S_a^{A_2}(1000, 16) = \frac{6907.76}{1431.73} \approx 4.82$$

$$S_a^{A_2}(1000, 64) = \frac{6907.76}{1107.93} \approx 6.23$$

3.2 Parallel Efficiency

One can calculate the parallel efficiency by computing $E^{A_i}(n, p) = \frac{S_a^{A_i}(n, p)}{p}$.

Algorithm A1:

$$E^{A_1}(1000, 4) = \frac{S_a^{A_1}(1000, 4)}{4} \approx 1.0$$

$$E^{A_1}(1000, 16) = \frac{S_a^{A_1}(1000, 16)}{16} \approx 0.98$$

$$E^{A_1}(1000, 64) = \frac{S_a^{A_1}(1000, 64)}{64} \approx 0.94$$

Algorithm A2:

$$E^{A_2}(1000, 4) = \frac{S_a^{A_2}(1000, 4)}{4} \approx 0.63$$

$$E^{A_2}(1000, 16) = \frac{S_a^{A_2}(1000, 16)}{16} \approx 0.30$$

$$E^{A_2}(1000, 64) = \frac{S_a^{A_2}(1000, 64)}{64} \approx 0.1$$

3.3 Potential Speed-up

According to Amdahl's Law the potential speed up is bounded by the sequential fraction of the code.

Algorithm A1:

The sequential fraction of A1 is $s_1 = \frac{\log n}{\frac{n \log n}{p} + \log n} = \frac{p}{n+p}$.

$$S^{A_1}(n, p) = \frac{1}{s_1 + \frac{1-s_1}{p}} \leq \frac{1}{s_1} = \frac{1}{\frac{p}{n+p} + \frac{1-\frac{p}{n+p}}{p}} \leq \frac{1}{\frac{p}{n+p}}$$

Algorithm A2:

The sequential fraction of A2 is $s_2 = \frac{\log n}{\frac{n \log n}{p} + n} = \frac{\log n * p}{n * (\log n + p)}$.

$$S^{A_2}(n, p) = \frac{1}{s_2 + \frac{1-s_2}{p}} \leq \frac{1}{s_2} = \frac{1}{\frac{\log n * p}{n * (\log n + p)} + \frac{1 - \frac{\log n * p}{n * (\log n + p)}}{p}} \leq \frac{1}{\frac{\log n * p}{n * (\log n + p)}}$$

However, in the solution above "p" is still involved. Therefore, we decided to use a different approach and take a look at the ratio of the sequential part of the algorithms (in terms of O notation) and take the ratio to the best performing sequential algorithm.

Algorithm A1:

The sequential part of A1 is $s_1 = \log n$.

$$S^{A_1}(n) = \frac{n \log n}{\log n} = n$$

Algorithm A2:

The sequential fraction of A2 is $s_2 = n$.

$$S^{A_2}(n) = \frac{n \log n}{n} = \log n$$

4 Weak Scaling Analysis

4.1 Growth of input size

$$O\left(\frac{n^4}{p}\right) = O(w)$$

$$\frac{n^4}{p} = w$$

$$n^4 = pw$$

$$n = (pw)^{\frac{1}{4}}$$

Since the average work per processor is w and the parallel algorithm runs in $O\left(\frac{n^4}{p}\right)$, n needs to increase as the fourth root of the product of w and p to keep the average work at $O(n)$.

4.2 Input size for different p

Given that $n=100$ when $p=1$, the various input sizes for different numbers of processors can be computed as follows. This time we decided to perform the calculations in R closely following the code presented in the lecture, as can be seen in listing 1.

Listing 1: Calculation n for various values of p

```
n <- 100
w <- n^4
df <- data.frame(p=c(2,4,8,16,128))
df$n <- ceiling((df$p*w)^(1/4))
print(df)
```

Table 4.2 shows the output of the code and presents the necessary input size for various numbers of cores.

p	n
2	119
4	142
8	169
16	200
128	337

Table 5: Required input sizes for different numbers of processors

Appendix

System Information

Listing 2: Output of 'lscpu' command

```
bopc23s9@hydra-head:~$ lscpu

Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
Address sizes:          46 bits physical, 48 bits virtual
CPU(s):                16
On-line CPU(s) list:   0-15
Thread(s) per core:    1
Core(s) per socket:    16
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:             6
Model:                  85
Model name:             Intel(R) Xeon(R) Gold 6130 CPU @ 2.10GHz
Stepping:               4
CPU MHz:                1000.151
CPU max MHz:            3700.0000
CPU min MHz:            1000.0000
BogoMIPS:               4200.00
L1d cache:              512 KiB
L1i cache:              512 KiB
L2 cache:               16 MiB
L3 cache:               22 MiB
NUMA node0 CPU(s):     0-15
```

Raw Data

size	patch	nprocs	time
155	26	1	0.27615245105698705
155	26	1	0.27596693905070424
155	26	1	0.2771365772932768
1100	26	1	13.435972596984357
1100	26	1	13.452274681068957
1100	26	1	13.371691829990596
155	26	2	0.14638189878314734
155	26	2	0.14576987363398075
155	26	2	0.1462678317911923
1100	26	2	6.670089309103787
1100	26	2	6.613465172238648
1100	26	2	6.7021133550442755
155	26	4	0.08490922581404448
155	26	4	0.08491231314837933
155	26	4	0.08519301796332002
1100	26	4	3.3561069620773196
1100	26	4	3.3553116391412914
1100	26	4	3.421122542116791
155	26	8	0.06856991816312075
155	26	8	0.06856326712295413
155	26	8	0.06828531296923757
1100	26	8	1.6898975907824934
1100	26	8	1.7937131081707776
1100	26	8	1.6963165369816124
155	26	16	0.07205457100644708
155	26	16	0.0730455950833857
155	26	16	0.07325727492570877
1100	26	16	0.9077705508098006
1100	26	16	0.9078279393725097
1100	26	16	0.9103577299974859
155	26	24	0.0845845378935337
155	26	24	0.08471266506239772
155	26	24	0.08427527407184243
1100	26	24	0.656599803827703
1100	26	24	0.6573387430980802
1100	26	24	0.6497617312707007
155	26	32	0.0973281990736723
155	26	32	0.09837603475898504
155	26	32	0.09687442006543279
1100	26	32	0.5734831769950688
1100	26	32	0.5541784320957959
1100	26	32	0.565086467191577

Table 6: Output 2.2.1

size	patch	nprocs	time
155	26	1	0.392238802742213
155	26	1	0.39877972193062305
155	26	1	0.3912382051348686
1100	26	1	19.34426457900554
1100	26	1	18.90565126761794
1100	26	1	19.064889647066593
155	26	2	0.20706806099042296
155	26	2	0.20968409720808268
155	26	2	0.2159389308653772
1100	26	2	9.65854894882068
1100	26	2	9.758664200082421
1100	26	2	9.597689433023334
155	26	4	0.12078673578798771
155	26	4	0.12187985517084599
155	26	4	0.12278737686574459
1100	26	4	4.8538289330899715
1100	26	4	4.949722504708916
1100	26	4	4.896979348734021
155	26	8	0.08609647722914815
155	26	8	0.08687825128436089
155	26	8	0.0856924457475543
1100	26	8	2.4588615149259567
1100	26	8	2.4403636367060244
1100	26	8	2.4585742950439453
155	26	16	0.08432547515258193
155	26	16	0.08534767106175423
155	26	16	0.08608504896983504
1100	26	16	1.2857805700041354
1100	26	16	1.28942183079198
1100	26	16	1.2837325409054756
155	26	24	0.09584450582042336
155	26	24	0.09703108435496688
155	26	24	0.09624077333137393
1100	26	24	0.9055881663225591
1100	26	24	0.8943913546390831
1100	26	24	0.9012130876071751
155	26	32	0.10889460984617472
155	26	32	0.10903614293783903
155	26	32	0.1076613599434495
1100	26	32	0.7539521278813481
1100	26	32	0.7489344119094312
1100	26	32	0.7355474359355867

Table 7: Output 2.2.2

size	patch	nprocs	time
950	1	32	60.38315498735756
950	1	32	58.993379454128444
950	1	32	59.664796941913664
950	5	32	2.4295669128187
950	5	32	2.457375079859048
950	5	32	2.430592040065676
950	10	32	0.8235339601524174
950	10	32	0.8456530389375985
950	10	32	0.8123667249456048
950	20	32	0.592178599908948
950	20	32	0.5885163000784814
950	20	32	0.5882865060120821
950	55	32	0.640842576045543
950	55	32	0.6622877516783774
950	55	32	0.6463527246378362
950	150	32	1.804339012131095
950	150	32	1.7855969751253724
950	150	32	1.7837719358503819
950	400	32	5.342996523249894
950	400	32	5.354430069215596
950	400	32	5.3442102540284395

Table 8: Output 2.3

size	patch	nprocs	time
800	1	16	5.177.733.509.801.320
800	1	16	48.679.752.216.208.700
800	1	16	49.604.969.315.230.800
800	2	16	12.208.631.441.928.400
800	2	16	12.400.218.938.011.600
800	2	16	1.267.593.849.496.910
800	3	16	5.231.156.553.607.430
800	3	16	5.173.088.782.932.600
800	3	16	530.134.785.734.117
800	4	16	2.795.454.104.896.630
800	4	16	28.061.751.988.716.400
800	4	16	2.743.914.455.641.060
800	5	16	17.844.559.531.658.800
800	5	16	1.795.137.454.289.940
800	5	16	1.754.671.474.918.720
800	6	16	12.681.594.076.566.300
800	6	16	12.242.525.001.056.400
800	6	16	12.942.032.269.202.100
800	7	16	10.078.809.931.874.200
800	7	16	10.316.235.939.972.100
800	7	16	10.379.106.737.673.200
800	8	16	0.8636039649136364
800	8	16	0.8836448309011757
800	8	16	0.8675666828639805
800	9	16	0.8217232823371887
800	9	16	0.8141451748088002
800	9	16	0.8879383755847812
800	10	16	0.8032627403736115
800	10	16	0.790965499356389
800	10	16	0.7719701924361289
800	11	16	0.775052934885025
800	11	16	0.7863136758096516
800	11	16	0.7699855691753328
800	12	16	0.7268759598955512
800	12	16	0.7416912410408258
800	12	16	0.7239395705983043
800	13	16	0.7274400270543993
800	13	16	0.7136987750418484
800	13	16	0.7242374992929399
800	14	16	0.7333642770536244
800	14	16	0.7298130737617612
800	14	16	0.7357385512441397
800	15	16	0.7113026860170066
800	15	16	0.7115092552267015

Table 9: Output 2.4 Part 1

size	patch	nprocs	time
800	15	16	0.7147729960270226
800	16	16	0.7343269921839237
800	16	16	0.7277633068151772
800	16	16	0.7470175549387932
800	17	16	0.7087195003405213
800	17	16	0.7273335340432823
800	17	16	0.7066979217343032
800	18	16	0.7008385430090129
800	18	16	0.7211737250909209
800	18	16	0.7053795196115971
800	19	16	0.704965204000473
800	19	16	0.7185817346908152
800	19	16	0.7119517708197236
800	20	16	0.7045390028506517
800	20	16	0.7090244409628212
800	20	16	0.6993431178852916
800	21	16	0.7067825132980943
800	21	16	0.7152607310563326
800	21	16	0.7141714952886105
800	22	16	0.7289159940555692
800	22	16	0.7070880490355194
800	22	16	0.7090400322340429
800	23	16	0.7122855372726917
800	23	16	0.7262964779511094
800	23	16	0.7180132349021733
800	24	16	0.7030764939263463
800	24	16	0.7233439306728542
800	24	16	0.7046146122738719
800	25	16	0.7067729285918176
800	25	16	0.7068284871056676
800	25	16	0.7069635195657611
800	26	16	0.7092408840544522
800	26	16	0.7142819189466536
800	26	16	0.7060380070470273
800	27	16	0.7240377422422171
800	27	16	0.703698709141463
800	27	16	0.7189373909495771
800	28	16	0.7107441178523004
800	28	16	0.7102009649388492
800	28	16	0.7038759361021221
800	29	16	0.7044810829684138
800	29	16	0.7004124890081584
800	29	16	0.7011993718333542
800	30	16	0.7012511510401964
800	30	16	0.7043010322377086
800	30	16	0.7103770151734352

16
Table 10: Output 2.4 Part 2